

## WALMART BUSINESS CASE STUDY

### PROBLEM STATEMENT : Walmart Black Friday Purchase Analysis

**Objective:** Analyze customer purchase behavior during Black Friday to understand spending patterns across different segments.

**Gender Focus:** Determine whether women spend more than men per transaction.

**Other Segments:** Examine how marital status (married vs unmarried) and age groups influence purchase amounts.

**Methodology:** Use transactional data to compute average spending, confidence intervals, and segment-wise trends.

**Business Impact:** Provide actionable insights to help Walmart optimize marketing strategies, promotions, and inventory planning for better business decisions.

### MODULES' IMPORT

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
import warnings
```

### DATA IMPORT

```
!wget https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/293/original/walmart_data.csv?1641285094
--2025-09-29 11:47:43-- https://d2beiqkhq929f0.cloudfront.net/public_assets/assets/000/001/293/original/walmart_data.csv?1641285094
Resolving d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)... 108.157.172.173, 108.157.172.183, 108.157.172.176, ...
Connecting to d2beiqkhq929f0.cloudfront.net (d2beiqkhq929f0.cloudfront.net)|108.157.172.173|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 23027994 (22M) [text/plain]
Saving to: 'walmart_data.csv?1641285094'

walmart_data.csv?16 100%[=====] 21.96M 141MB/s in 0.2s
2025-09-29 11:47:43 (141 MB/s) - 'walmart_data.csv?1641285094' saved [23027994/23027994]
```

```
df = pd.read_csv('walmart_data.csv?1641285094')
df.sample(5)
```

	User_ID	Product_ID	Gender	Age	Occupation	City_Category	Stay_In_Current_City_Years	Marital_Status	Product_Category	Purchase
310776	1005889	P00256642	M	51-55	20	C		1	1	5 5443
350998	1000033	P00273442	M	46-50	3	A		1	1	12 1395
242567	1001383	P00214142	F	26-35	7	A		1	0	8 5853
282537	1001519	P00116842	M	18-25	4	C		1	0	2 16498
407846	1002811	P0097742	M	26-35	19	B		1	0	8 5922

We observe a population of ~550k male and female customers. Considering this as a substantially large population, inferences from the samples are valid provided they are random and independent.

### EXPLORATORY DATA ANALYTICS

#### Data Overview

```
df.shape
```

```
(550068, 10)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 550068 entries, 0 to 550067
Data columns (total 10 columns):
 #   Column           Non-Null Count  Dtype  
 ---  -- 
 0   User_ID          550068 non-null   int64  
 1   Product_ID       550068 non-null   int64  
 2   Gender           550068 non-null   object 
 3   Age              550068 non-null   object 
 4   Occupation       550068 non-null   object 
 5   City_Category    550068 non-null   object 
 6   Stay_In_Current_City_Years 550068 non-null   int64  
 7   Marital_Status   550068 non-null   object 
 8   Product_Category 550068 non-null   int64  
 9   Purchase         550068 non-null   int64 
```

```

0 User_ID           550068 non-null int64
1 Product_ID        550068 non-null object
2 Gender            550068 non-null object
3 Age               550068 non-null object
4 Occupation        550068 non-null int64
5 City_Category     550068 non-null object
6 Stay_In_Current_City_Years 550068 non-null object
7 Marital_Status    550068 non-null int64
8 Product_Category  550068 non-null int64
9 Purchase          550068 non-null int64
dtypes: int64(5), object(5)
memory usage: 42.0+ MB

```

```
df.isnull().sum() # NO NULL VALUES
```

	0
User_ID	0
Product_ID	0
Gender	0
Age	0
Occupation	0
City_Category	0
Stay_In_Current_City_Years	0
Marital_Status	0
Product_Category	0
Purchase	0

```
dtype: int64
```

```
df[df.duplicated()].shape # NO DUPLICATE ROWS
```

```
(0, 10)
```

```

df['Gender'].unique(), df['Occupation'].unique(), df['City_Category'].unique(), df['Marital_Status'].unique(), df['Product_Category'].unique()

(array(['F', 'M'], dtype=object),
 array([10, 16, 15, 7, 20, 9, 1, 12, 17, 0, 3, 4, 11, 8, 19, 2, 18,
       5, 14, 13, 6]),
 array(['A', 'C', 'B'], dtype=object),
 array([0, 1]),
 array([ 3,  1, 12,  8,  5,  4,  2,  6, 14, 11, 13, 15,  7, 16, 18, 10, 17,
        9, 20, 19]))

```

### Converting categories to category dtype:

```

df['Gender'] = df['Gender'].astype('category')
df['Occupation'] = df['Occupation'].astype('category')
df['City_Category'] = df['City_Category'].astype('category')
df['Marital_Status'] = df['Marital_Status'].astype('category')
df['Product_Category'] = df['Product_Category'].astype('category')

```

```

print("Gender codes: ", df["Gender"].cat.codes.unique())
print("Occupation codes: ", df["Occupation"].cat.codes.unique())
print("City_Category codes: ", df["City_Category"].cat.codes.unique())
print("Marital_Status codes: ", df["Marital_Status"].cat.codes.unique())
print("Product_Category codes: ", df["Product_Category"].cat.codes.unique())

```

```

Gender codes: [0 1]
Occupation codes: [10 16 15 7 20 9 1 12 17 0 3 4 11 8 19 2 18 5 14 13 6]
City_Category codes: [0 2 1]
Marital_Status codes: [0 1]
Product_Category codes: [ 2  0 11  7  4  3  1  5 13 10 12 14  6 15 17  9 16  8 19 18]

```

### Converting Age to ordered categorical dtype: Sorting will follow life stages, not alphabet (string values)

```

print(df['Age'].unique())
df['Age'].dtype

['0-17' '55+' '26-35' '46-50' '51-55' '36-45' '18-25']
dtype('O')

```

```

age_order = ["0-17", "18-25", "26-35", "36-45", "46-50", "51-55", "55+"]
df['Age'] = pd.Categorical(df['Age'], categories=age_order, ordered=True)

```

```

print("Age codes: ", df["Age"].cat.codes.unique())
Age codes: [0 6 2 4 5 3 1]

```

## Statistical Summary

```
df.describe().T
```

	count	mean	std	min	25%	50%	75%	max
User_ID	550068.0	1.003029e+06	1727.591586	1000001.0	1001516.0	1003077.0	1004478.0	1006040.0
Purchase	550068.0	9.263969e+03	5023.065394	12.0	5823.0	8047.0	12054.0	23961.0

```
cat_cols = ['Gender', 'Age', 'City_Category', 'Marital_Status']
df[cat_cols].melt().groupby(['variable', 'value'])[['value']].count()/len(df)*100
```

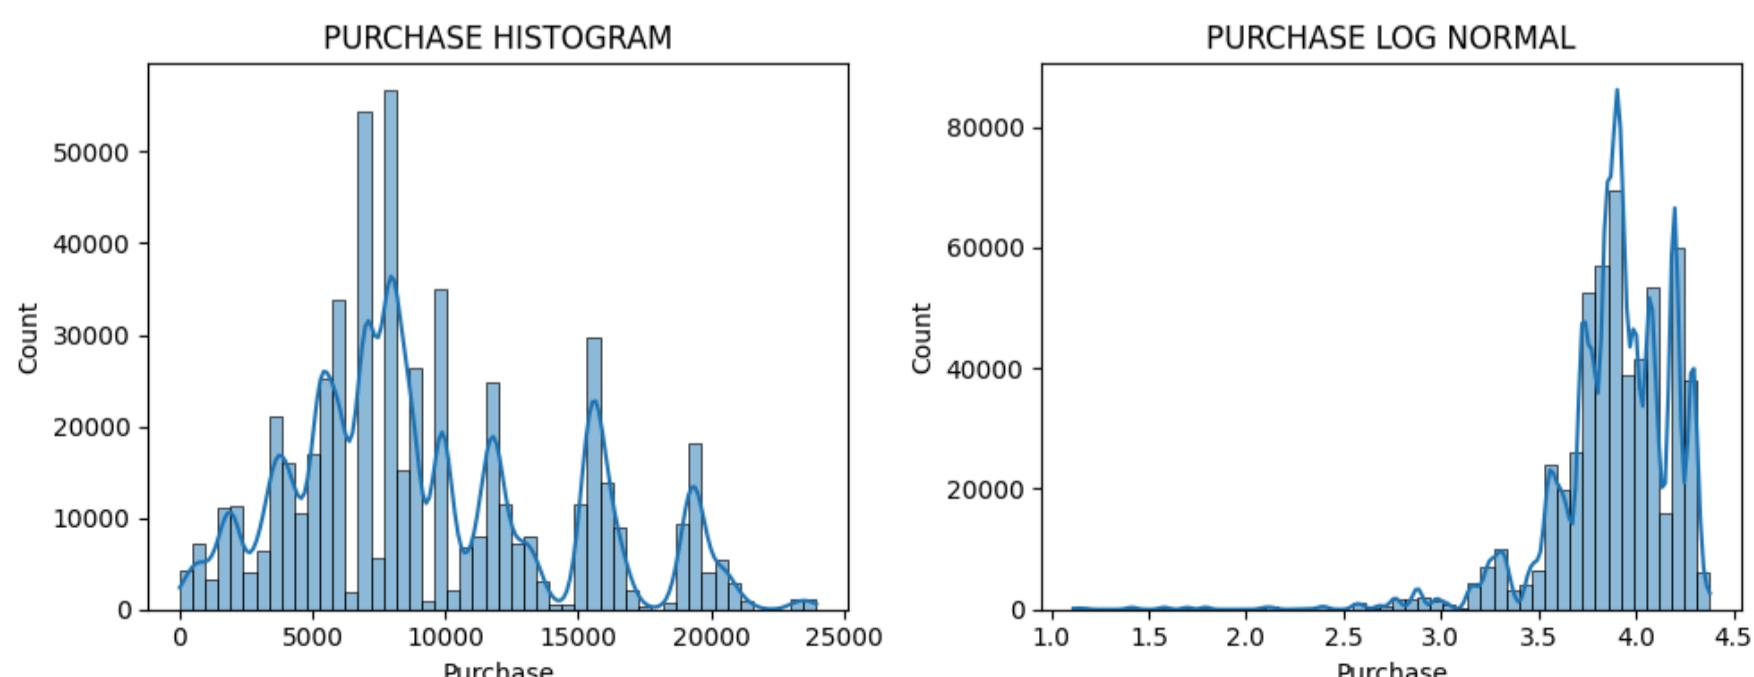
	variable	value	
Age	0-17	2.745479	
	18-25	18.117760	
	26-35	39.919974	
	36-45	19.999891	
	46-50	8.308246	
	51-55	6.999316	
	55+	3.909335	
City_Category	A	26.854862	
	B	42.026259	
	C	31.118880	
Gender	F	24.689493	
	M	75.310507	
Marital_Status	0	59.034701	
	1	40.965299	

- 60% of the users fall within the age bracket of 26–45 years, which shows that the majority of Walmart's customers are young to middle-aged adults.
- 42% of the users belong to City Category B, suggesting that this city type contributes the largest share of customers.
- 75% of the users are male, indicating that the customer base is heavily dominated by men.
- 60% of the users are single, highlighting that unmarried individuals form the majority of Walmart's customer base.

## Visualizing Univariate Distributions

### Visualizing Continuous Variables

```
fig, axis = plt.subplots(1, 2, figsize=(10, 4))
sns.histplot(data=df, x='Purchase', ax=axis[0], bins=50, kde=True)
axis[0].set_title('PURCHASE HISTOGRAM')
sns.histplot(np.log10(df['Purchase']+1), ax=axis[1], bins=50, kde=True)
axis[1].set_title('PURCHASE LOG NORMAL')
plt.tight_layout()
plt.show()
```

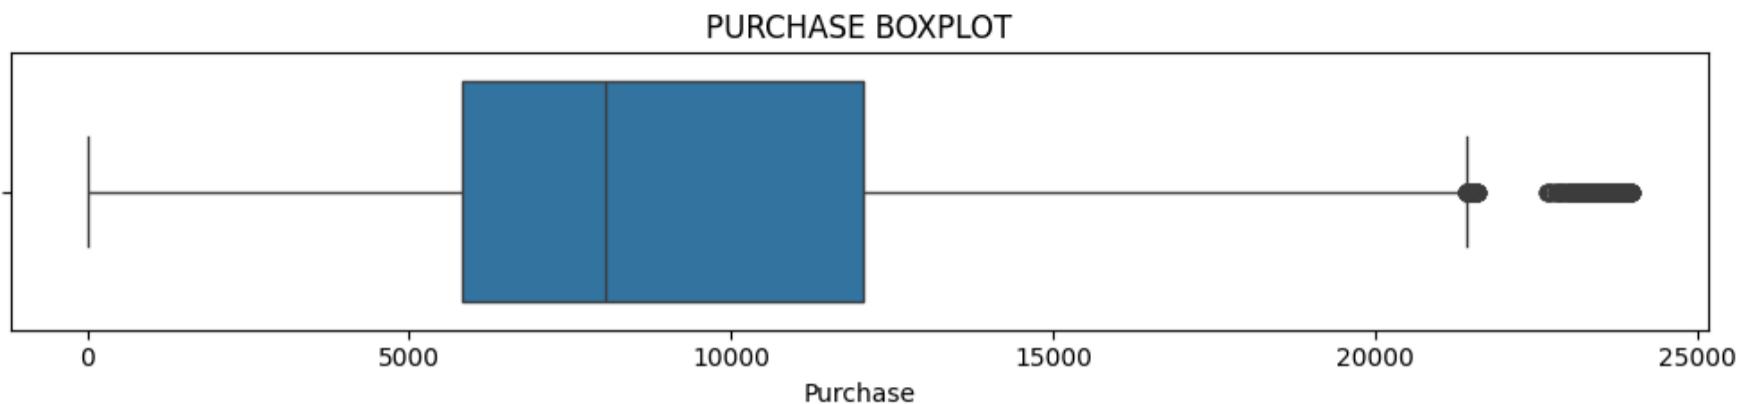


```
fig, axis = plt.subplots(1, 1, figsize=(10, 2.5))
sns.boxplot(x=df['Purchase'])
```

```

plt.title('PURCHASE BOXPLOT')
plt.tight_layout()
plt.show()

```



```

q1 = df['Purchase'].quantile(0.25)
q3 = df['Purchase'].quantile(0.75)
iqr = q3 - q1

lower_bound = q1 - 1.5 * iqr
upper_bound = q3 + 1.5 * iqr

outliers = df[(df['Purchase'] < lower_bound) | (df['Purchase'] > upper_bound)]

print("Number of outliers:", len(outliers))
print("Percentage of outliers:", (len(outliers)/len(df))*100)

```

Number of outliers: 2677  
Percentage of outliers: 0.4866671029763593

## ▼ Visualizing Categorical Variables

```

df.columns

```

```

Index(['User_ID', 'Product_ID', 'Gender', 'Age', 'Occupation', 'City_Category',
       'Stay_In_Current_City_Years', 'Marital_Status', 'Product_Category',
       'Purchase'],
      dtype='object')

```

```

fig, axis = plt.subplots(2,2, figsize=(10,5))

```

```

sns.countplot(data=df, x='Gender', ax=axis[0,0], color='lightgrey', edgecolor='black')

```

```

axis[0,0].set_title('GENDER')

```

```

sns.countplot(data=df, x='Age', ax=axis[0,1], color='lightgrey', edgecolor='black')

```

```

axis[0,1].set_title('AGE')

```

```

sns.countplot(data=df, x='Marital_Status', ax=axis[1,0], color='lightgrey', edgecolor='black')

```

```

axis[1,0].set_title('MARITAL STATUS')

```

```

sns.countplot(data=df, x='City_Category', ax=axis[1,1], color='lightgrey', edgecolor='black')

```

```

axis[1,1].set_title('CITY CATEGORY')

```

```

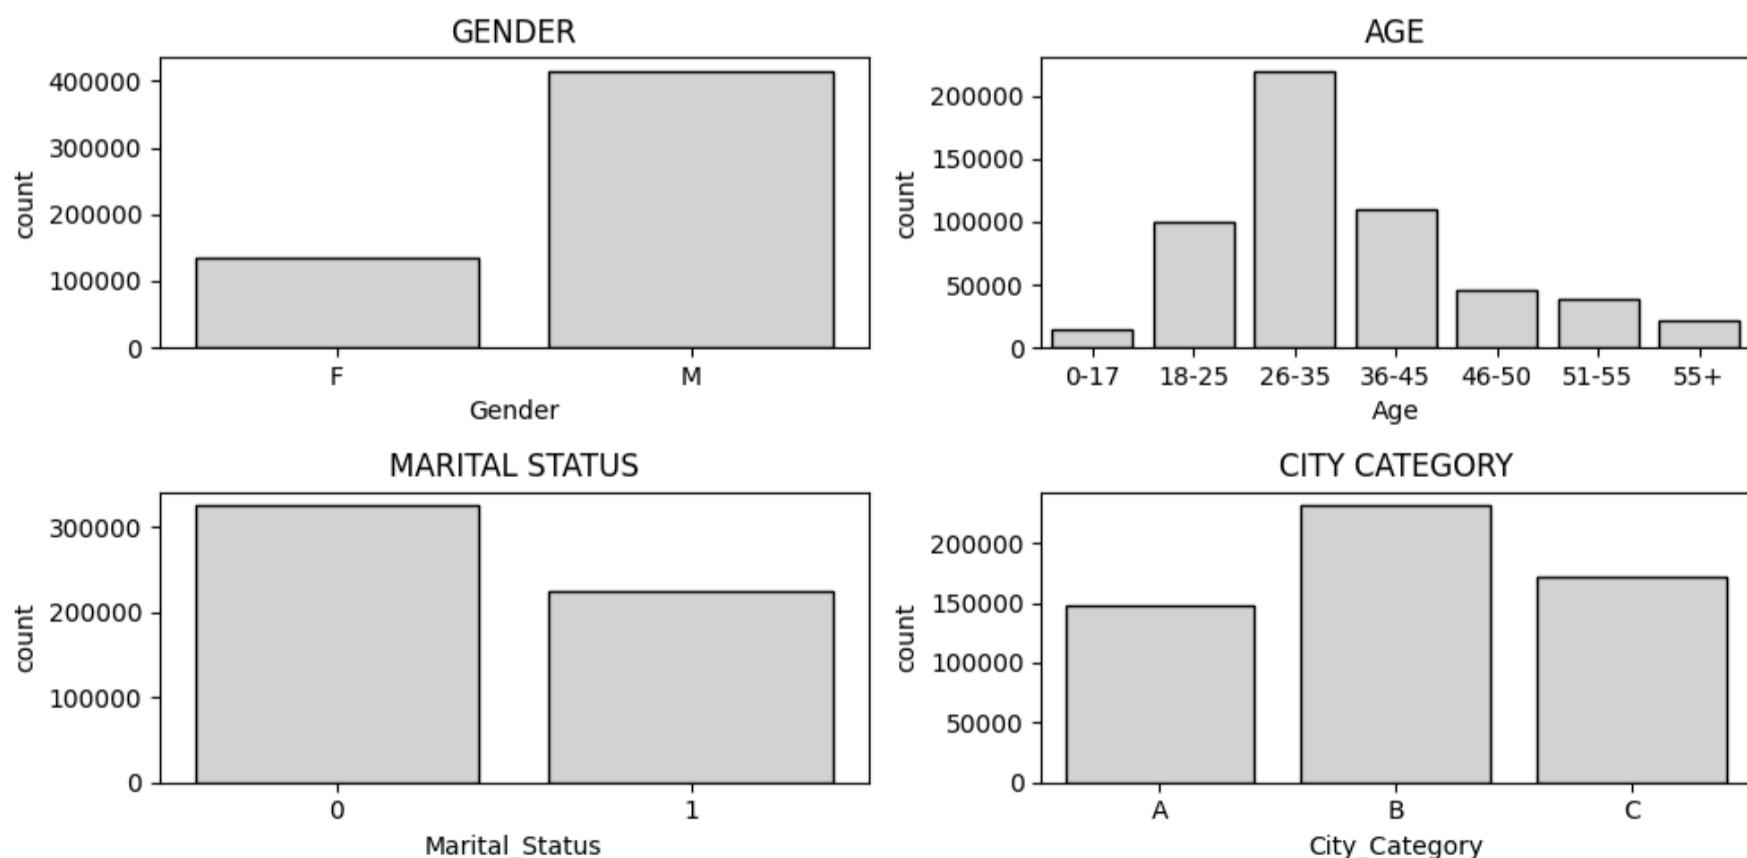
plt.tight_layout()

```

```

plt.show()

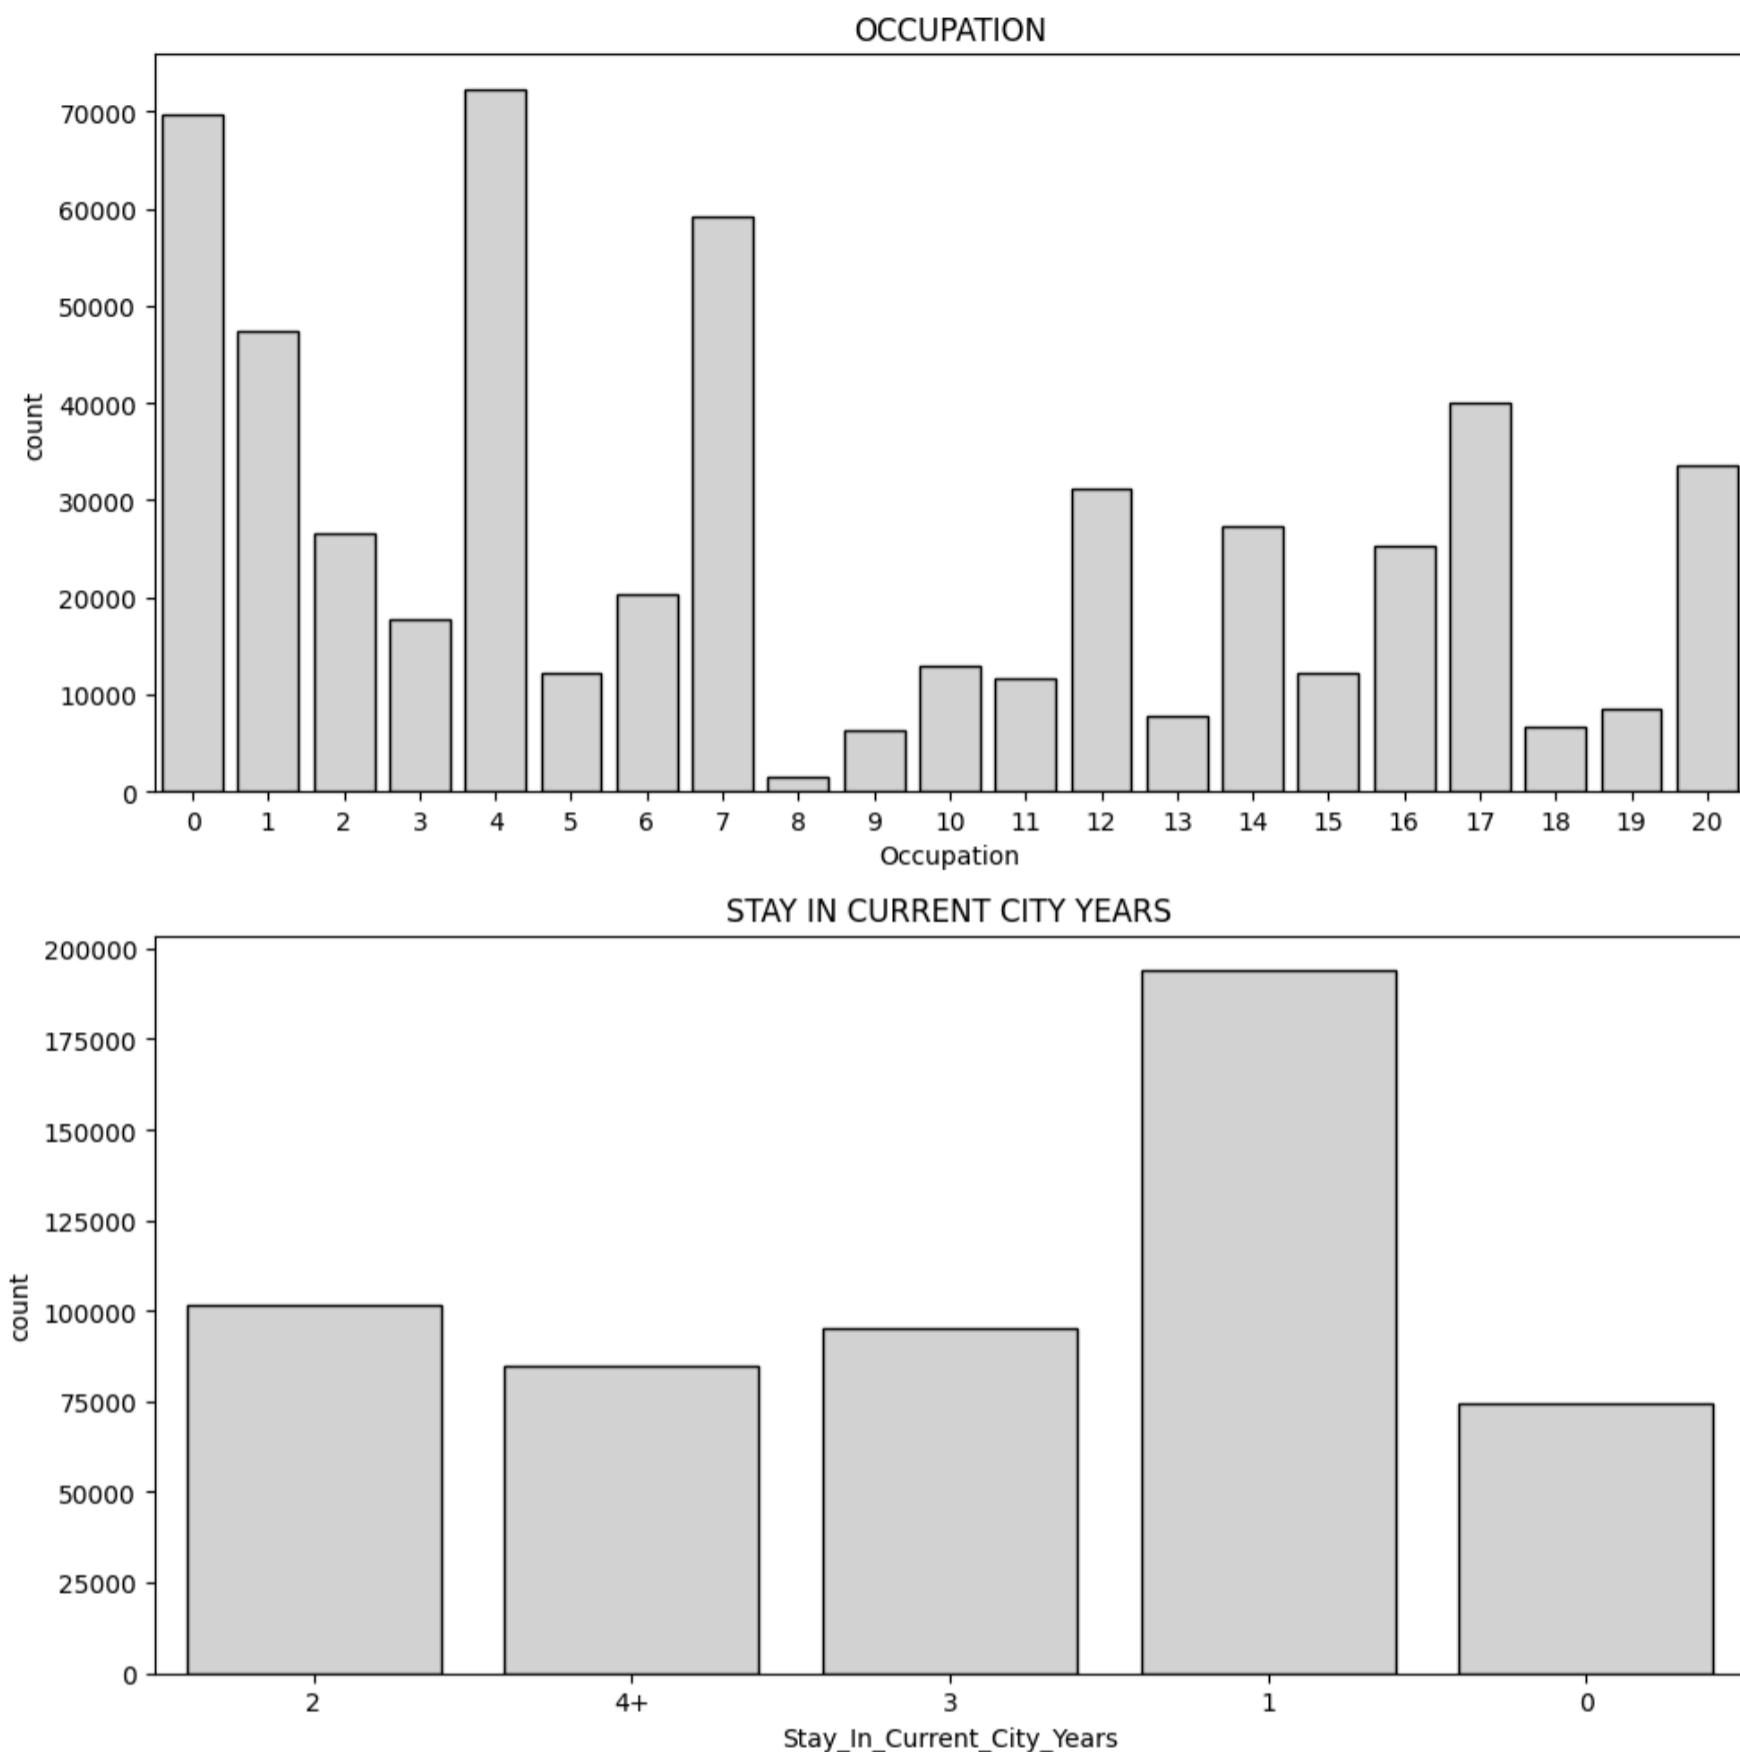
```



```

fig, axis = plt.subplots(2,1, figsize=(10,10))
sns.countplot(data=df, x='Occupation', ax=axis[0], color='lightgrey', edgecolor='black')
axis[0].set_title('OCCUPATION')
sns.countplot(data=df, x='Stay_In_Current_City_Years', ax=axis[1], color='lightgrey', edgecolor='black')
axis[1].set_title('STAY IN CURRENT CITY YEARS')
plt.tight_layout()
plt.show()

```



## ▼ Bivariate Analysis

---

## ▼ Purchase vis-a-vis Gender

---

```

df.groupby('Gender')[ 'Purchase'].describe()

/tmppython-input-1243678314.py:1: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version
df.groupby('Gender')[ 'Purchase'].describe()

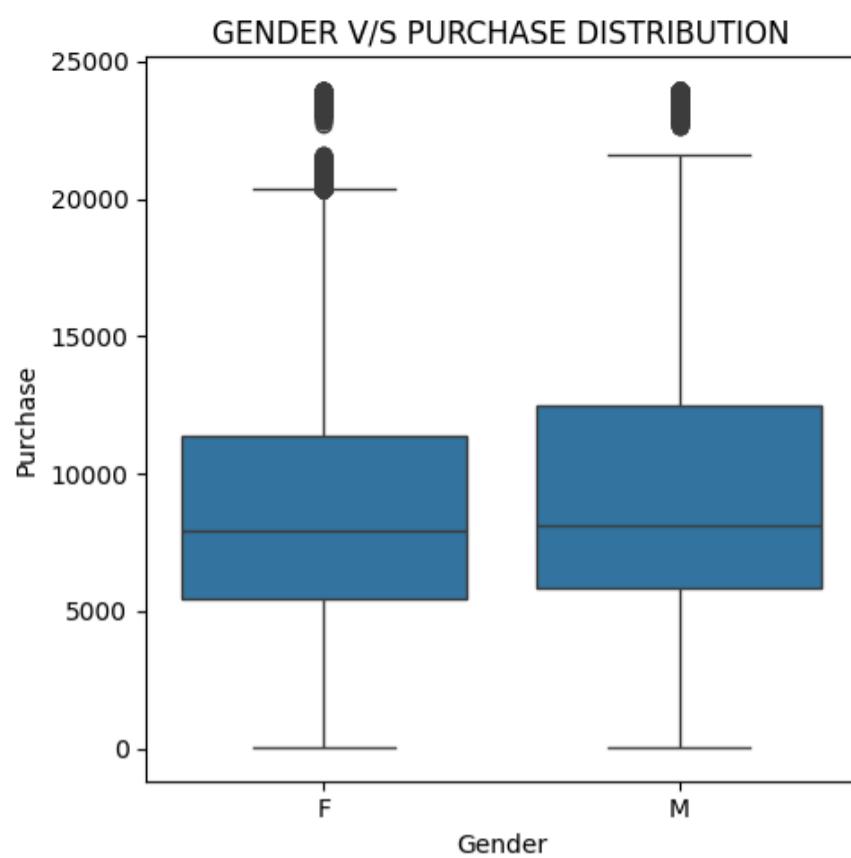
      count        mean         std       min     25%     50%     75%       max
Gender
F    135809.0  8734.565765  4767.233289   12.0  5433.0  7914.0  11400.0  23959.0
M    414259.0  9437.526040  5092.186210   12.0  5863.0  8098.0  12454.0  23961.0

```

```

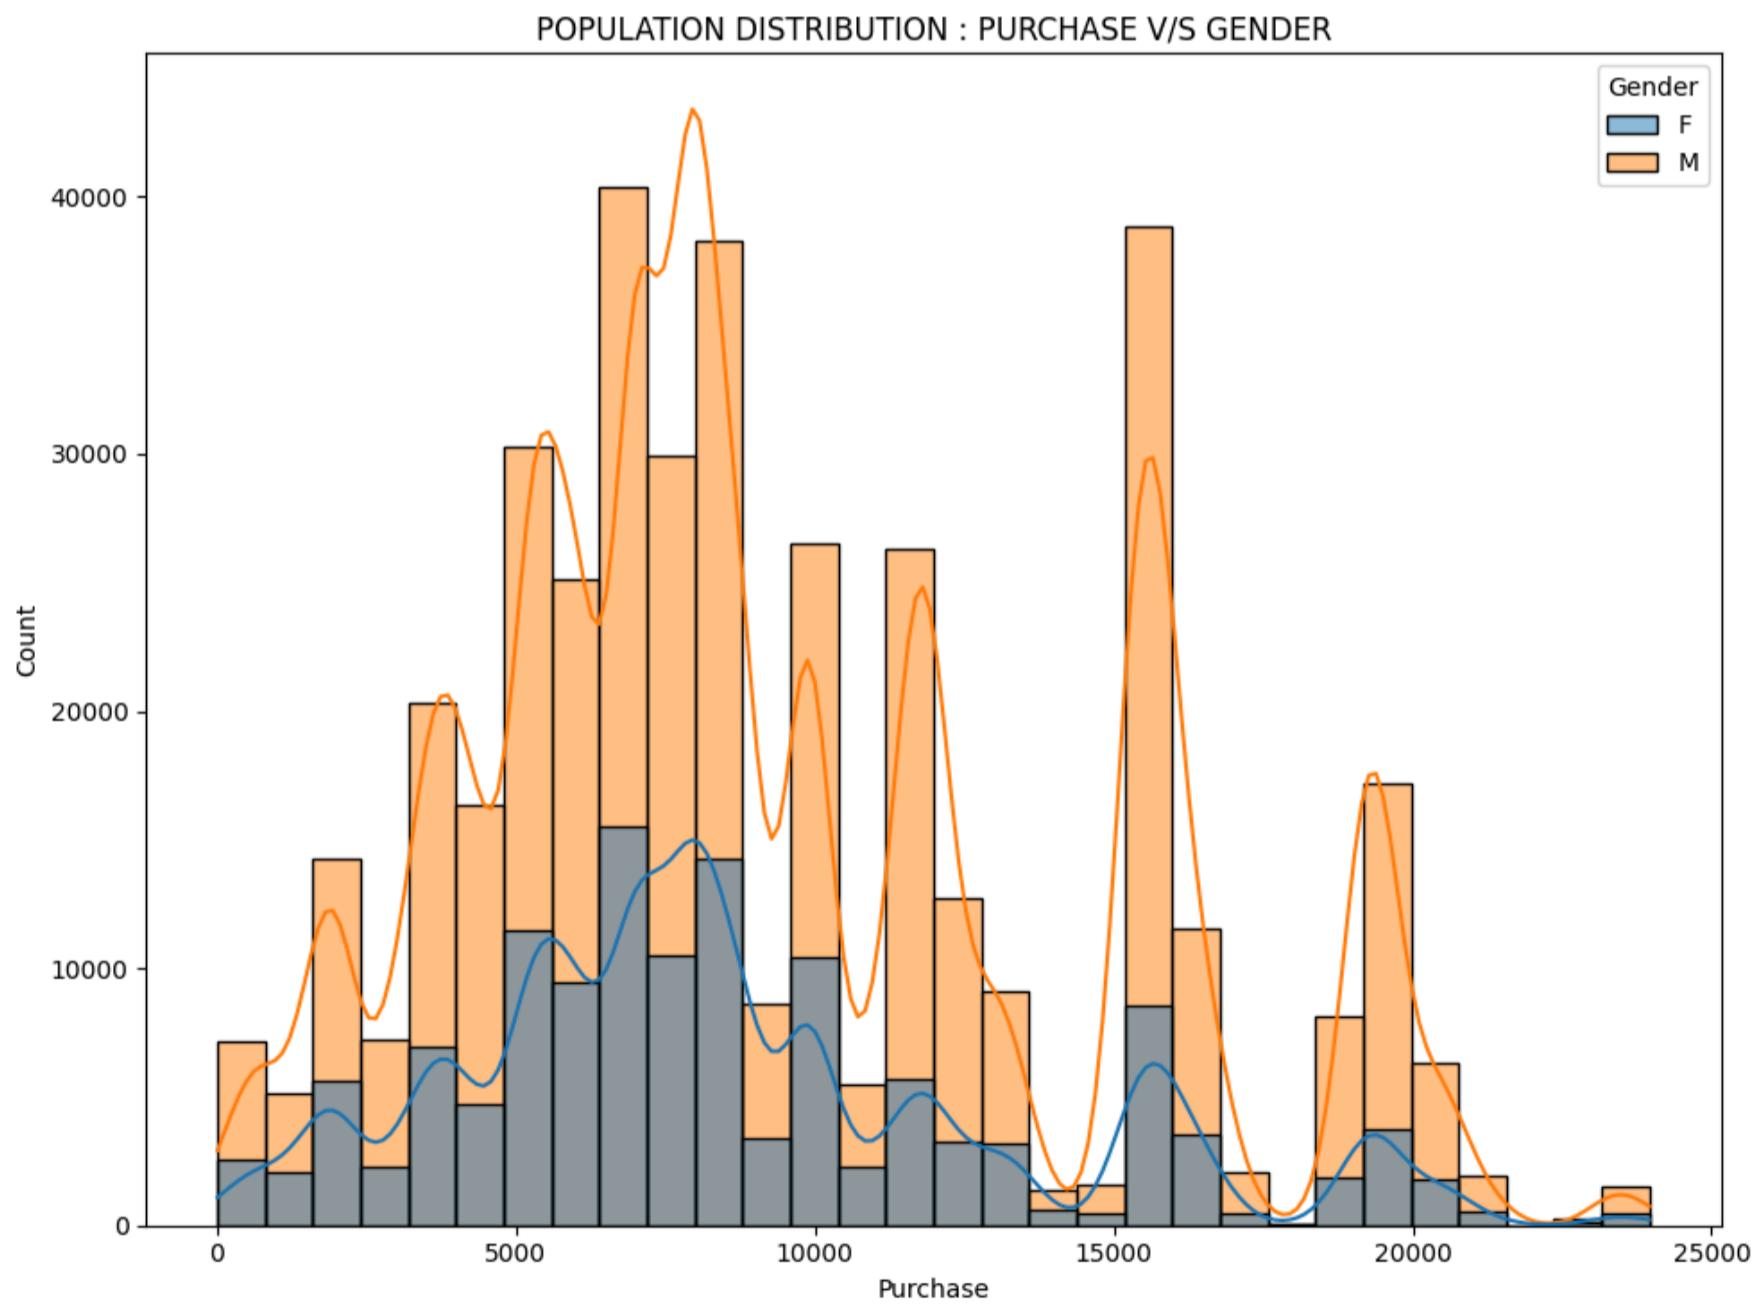
fig, axis = plt.subplots(1,1, figsize=(5,5))
sns.boxplot(data=df, x='Gender', y='Purchase', ax=axis)
axis.set_title('GENDER V/S PURCHASE DISTRIBUTION')
plt.tight_layout()
plt.show()

```



### Population Distribution Analysis

```
fig, axis = plt.subplots(1,1, figsize=(10,7.5))
sns.histplot(data=df, x='Purchase', bins=30, kde=True, hue='Gender', ax=axis)
plt.title('POPULATION DISTRIBUTION : PURCHASE V/S GENDER')
plt.tight_layout()
plt.show()
```



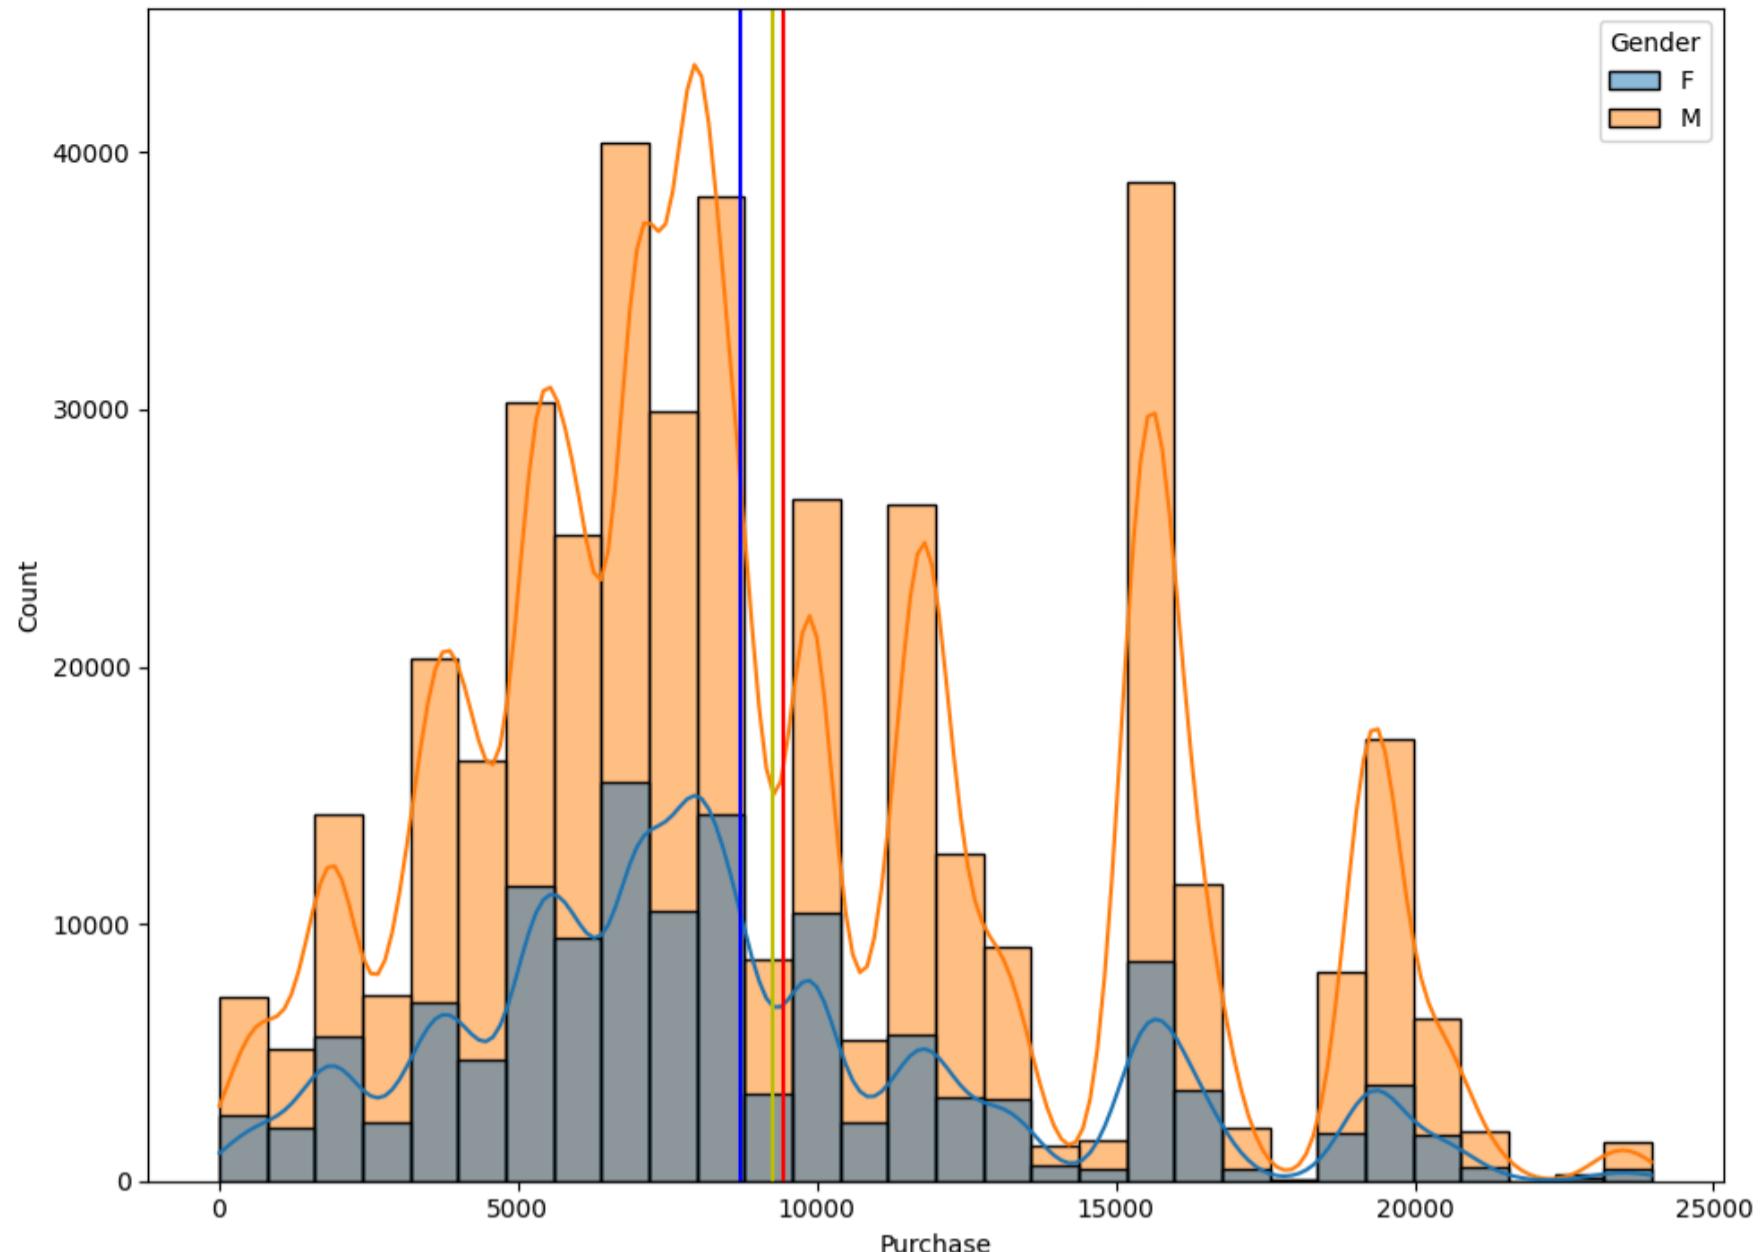
Notes: The distribution appears roughly normal, though an outlier is evident around 15,000

### Population Mean Analysis

```
fig, axis = plt.subplots(1,1, figsize=(10,7.5))
sns.histplot(data=df, x='Purchase', bins=30, kde=True, hue='Gender', ax=axis)
plt.axvline(x=df['Purchase'].mean(), color='y')
plt.axvline(x=df[df['Gender']=='M']['Purchase'].mean(), color='r')
plt.axvline(x=df[df['Gender']=='F']['Purchase'].mean(), color='b')
plt.title('POPULATION MEAN : MALE V/S FEMALE V/S COMBINED')
```

```
plt.tight_layout()
plt.show()
```

POPULATION MEAN : MALE V/S FEMALE V/S COMBINED



Notes: The mean spending for males, females, and the overall sample is similar, with females spending slightly less

#### SAMPLE-DERIVED GENDER-CENTRIC ANALYSIS USING 1000 SAMPLE POINTS

##### FEMALES

```
df_f = df[df['Gender']=='F']
print("Mean of Purchase Amount by Females:",df_f['Purchase'].mean())
print("Standard Deviation of Purchase Amount by Females:",df_f['Purchase'].std())

Mean of Purchase Amount by Females: 8734.565765155476
Standard Deviation of Purchase Amount by Females: 4767.233289291444
```

```
sample_mean_of_purchase_by_females = []
for i in range(1000):
    sample_f = df_f['Purchase'].sample(n=7500)
    sample_f_mean = sample_f.mean()
    sample_mean_of_purchase_by_females.append(sample_f_mean)
sample_mean_of_purchase_by_females = np.array(sample_mean_of_purchase_by_females)

sample_mean_of_purchase_by_females[:10]

array([8775.50546667, 8680.09173333, 8838.23413333, 8733.1336,
       8767.11706667, 8769.03186667, 8776.6692, 8729.47586667,
       8778.35293333, 8807.84386667])
```

##### MALES

```
df_m = df[df['Gender']=='M']
print("Mean of Purchase Amount by Males:",df_m['Purchase'].mean())
print("Standard Deviation of Purchase Amount by Males:",df_m['Purchase'].std())

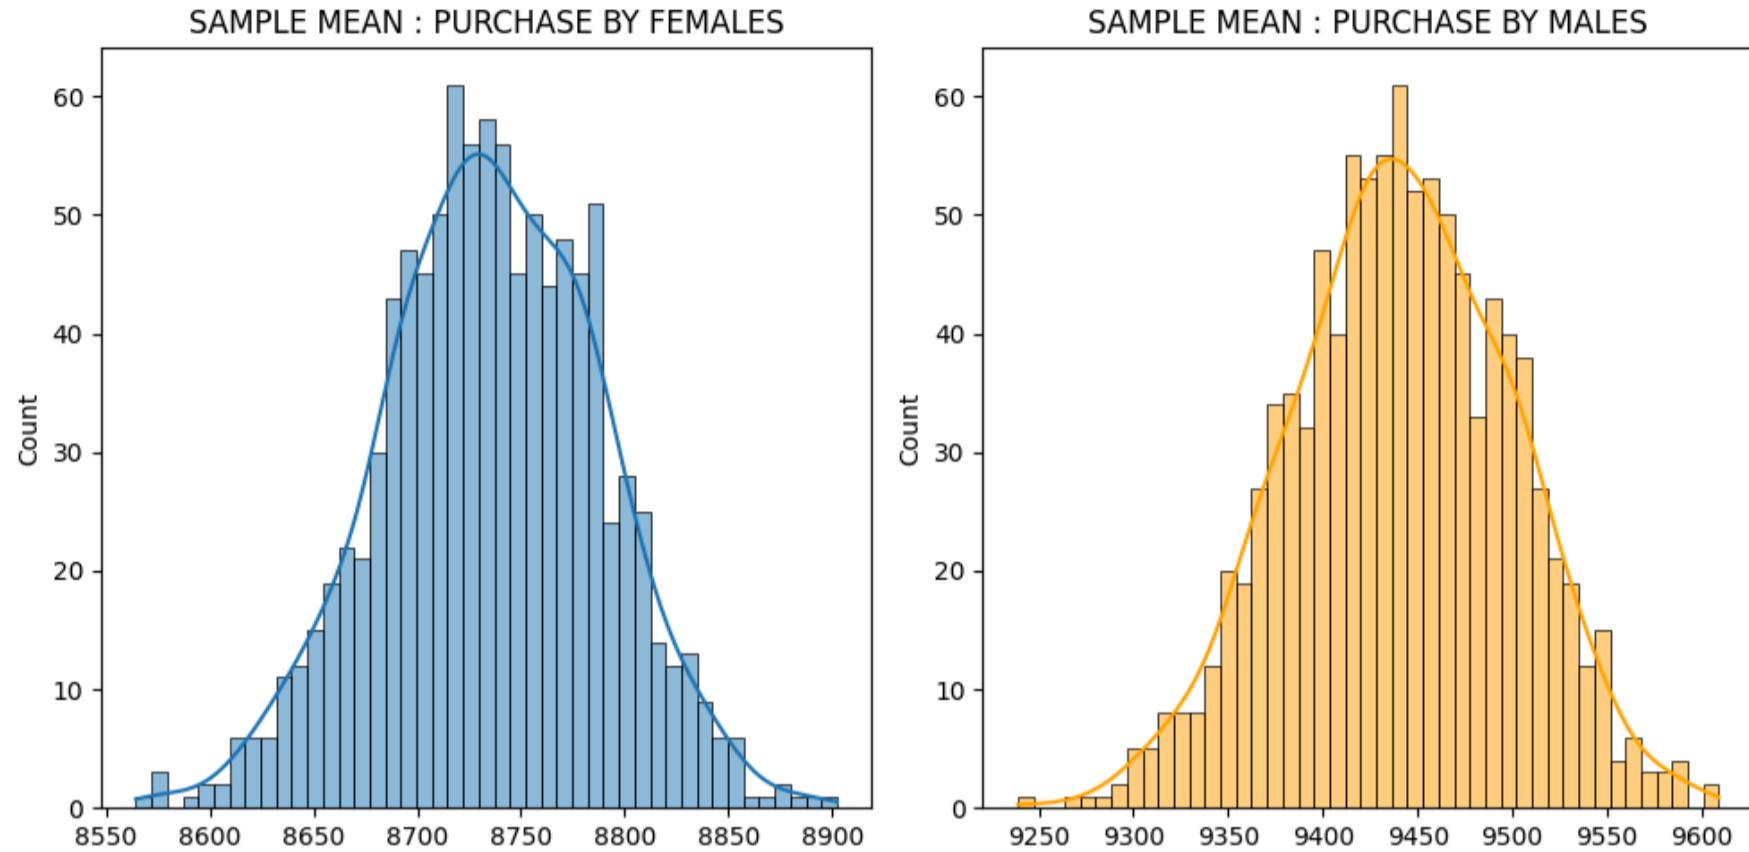
Mean of Purchase Amount by Males: 9437.526040472265
Standard Deviation of Purchase Amount by Males: 5092.186209777949
```

```
sample_mean_of_purchase_by_males = []
for i in range(1000):
    sample_m = df_m['Purchase'].sample(n=7500)
    sample_m_mean = sample_m.mean()
    sample_mean_of_purchase_by_males.append(sample_m_mean)
sample_mean_of_purchase_by_males = np.array(sample_mean_of_purchase_by_males)
```

```
sample_mean_of_purchase_by_males[:10]

array([9342.21493333, 9356.6844, 9512.2004, 9471.31453333,
       9380.058, 9461.36, 9456.962, 9415.81666667,
       9466.61493333, 9436.20933333])
```

```
fig, axis = plt.subplots(1, 2, figsize=(10, 5))
sns.histplot(data=sample_mean_of_purchase_by_females, bins=45, kde=True, ax=axis[0])
axis[0].set_title('SAMPLE MEAN : PURCHASE BY FEMALES')
sns.histplot(data=sample_mean_of_purchase_by_males, bins=45, kde=True, ax=axis[1], color='orange')
axis[1].set_title('SAMPLE MEAN : PURCHASE BY MALES')
plt.tight_layout()
plt.show()
```



### SAMPLE DISTRIBUTION MEANS:

```
print("Sample Mean of Purchase Amount by Females:", sample_mean_of_purchase_by_females.mean())
print("Sample Mean of Purchase Amount by Males:", sample_mean_of_purchase_by_males.mean())

Sample Mean of Purchase Amount by Females: 8735.427967733334
Sample Mean of Purchase Amount by Males: 9440.581635066666
```

The mean of this sampling distribution was found to be very close to the original population mean.

#### Females

- Population mean = 8734.56
- Sampling distribution mean = 8735.42

#### Males

- Population mean = 9437.52
- Sampling distribution mean = 9440.58

This justifies the first property of the Central Limit Theorem: Sampling Distribution Mean

≈

Population Mean

### ESTIMATION OF THE 95% CONFIDENCE INTERVAL

```
from scipy.stats import norm
```

#### FEMALES

```
mu = df_f['Purchase'].mean()
sigma = df_f['Purchase'].std()
n = len(df_f)
alpha = 0.05
se = sigma/np.sqrt(n)
z1 = norm.ppf(alpha/2)
z2 = norm.ppf(1-alpha/2)
x1 = z1*se + mu
x2 = z2*se + mu
print("Confidence Interval for Females:", (x1, x2))
```

```
Confidence Interval for Females: (np.float64(8709.21154714068), np.float64(8759.919983170272))
```

## MALES

```
mu = df_m['Purchase'].mean()
sigma = df_m['Purchase'].std()
n = len(df_m)
alpha = 0.05
se = sigma/np.sqrt(n)
z1 = norm.ppf(alpha/2)
z2 = norm.ppf(1-alpha/2)
x1 = z1*se + mu
x2 = z2*se + mu
print("Confidence Interval for Males:",(x1,x2))

Confidence Interval for Males: (np.float64(9422.01944736257), np.float64(9453.032633581959))
```

The **95% confidence intervals** in each case included the true population values — **consistent with the Central Limit Theorem**.

## Numerical column v/s categorical column with two unique entries => Independent Samples T-test

```
from scipy.stats import ttest_ind

# H0 : Purchase amount in independent of Gender
# H1 : Purchase amount in dependent on Gender

sample_size = 10000

male_sample = df[df['Gender'] == 'M']['Purchase'].sample(n=sample_size, random_state=42)
female_sample = df[df['Gender'] == 'F']['Purchase'].sample(n=sample_size, random_state=42)

result = ttest_ind(male_sample, female_sample, alternative='two-sided')
alpha = 0.05

if result.pvalue < alpha:
    print("Reject Null Hypothesis")
    print("Gender has a significant impact on Purchase Amount\n")
else:
    print("Fail to Reject Null Hypothesis")
    print("Gender has no significant impact on Purchase Amount\n")

result

Reject Null Hypothesis
Gender has a significant impact on Purchase Amount

TtestResult(statistic=np.float64(10.842139195163371), pvalue=np.float64(2.5889036301871065e-27), df=np.float64(19998.0))
```

## ▼ Purchase vis-a-vis Marital Status

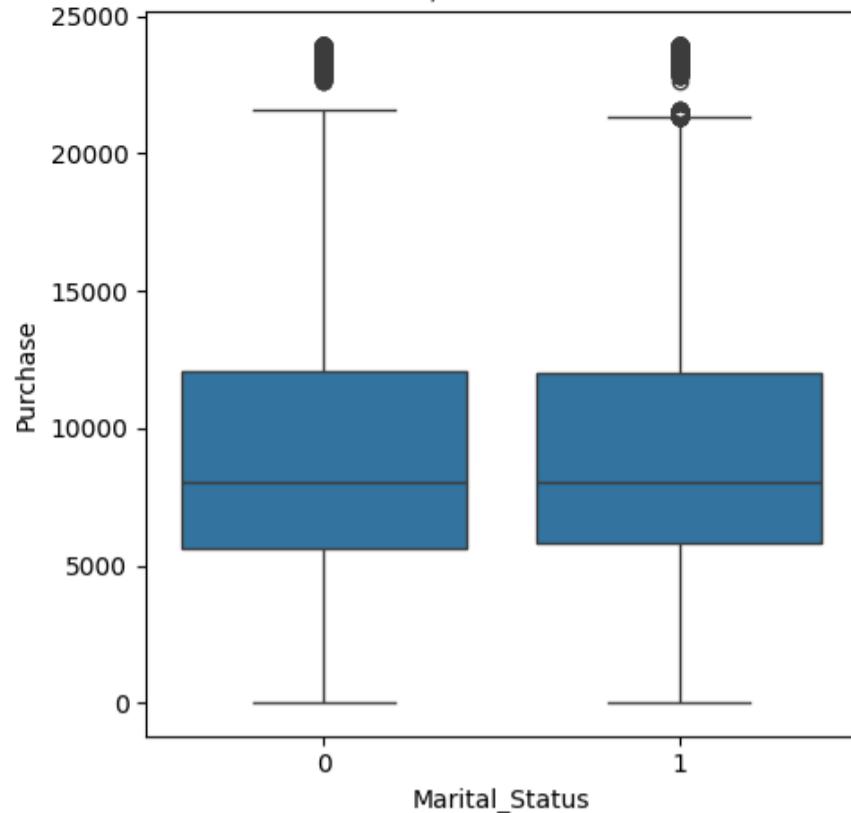
```
df.groupby('Marital_Status')['Purchase'].describe()

/tmp/ipython-input-3740306556.py:1: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version
df.groupby('Marital_Status')['Purchase'].describe()

      count        mean         std      min      25%      50%      75%      max
Marital_Status
  0    324731.0  9265.907619  5027.347859  12.0  5605.0  8044.0  12061.0  23961.0
  1    225337.0  9261.174574  5016.897378  12.0  5843.0  8051.0  12042.0  23961.0
```

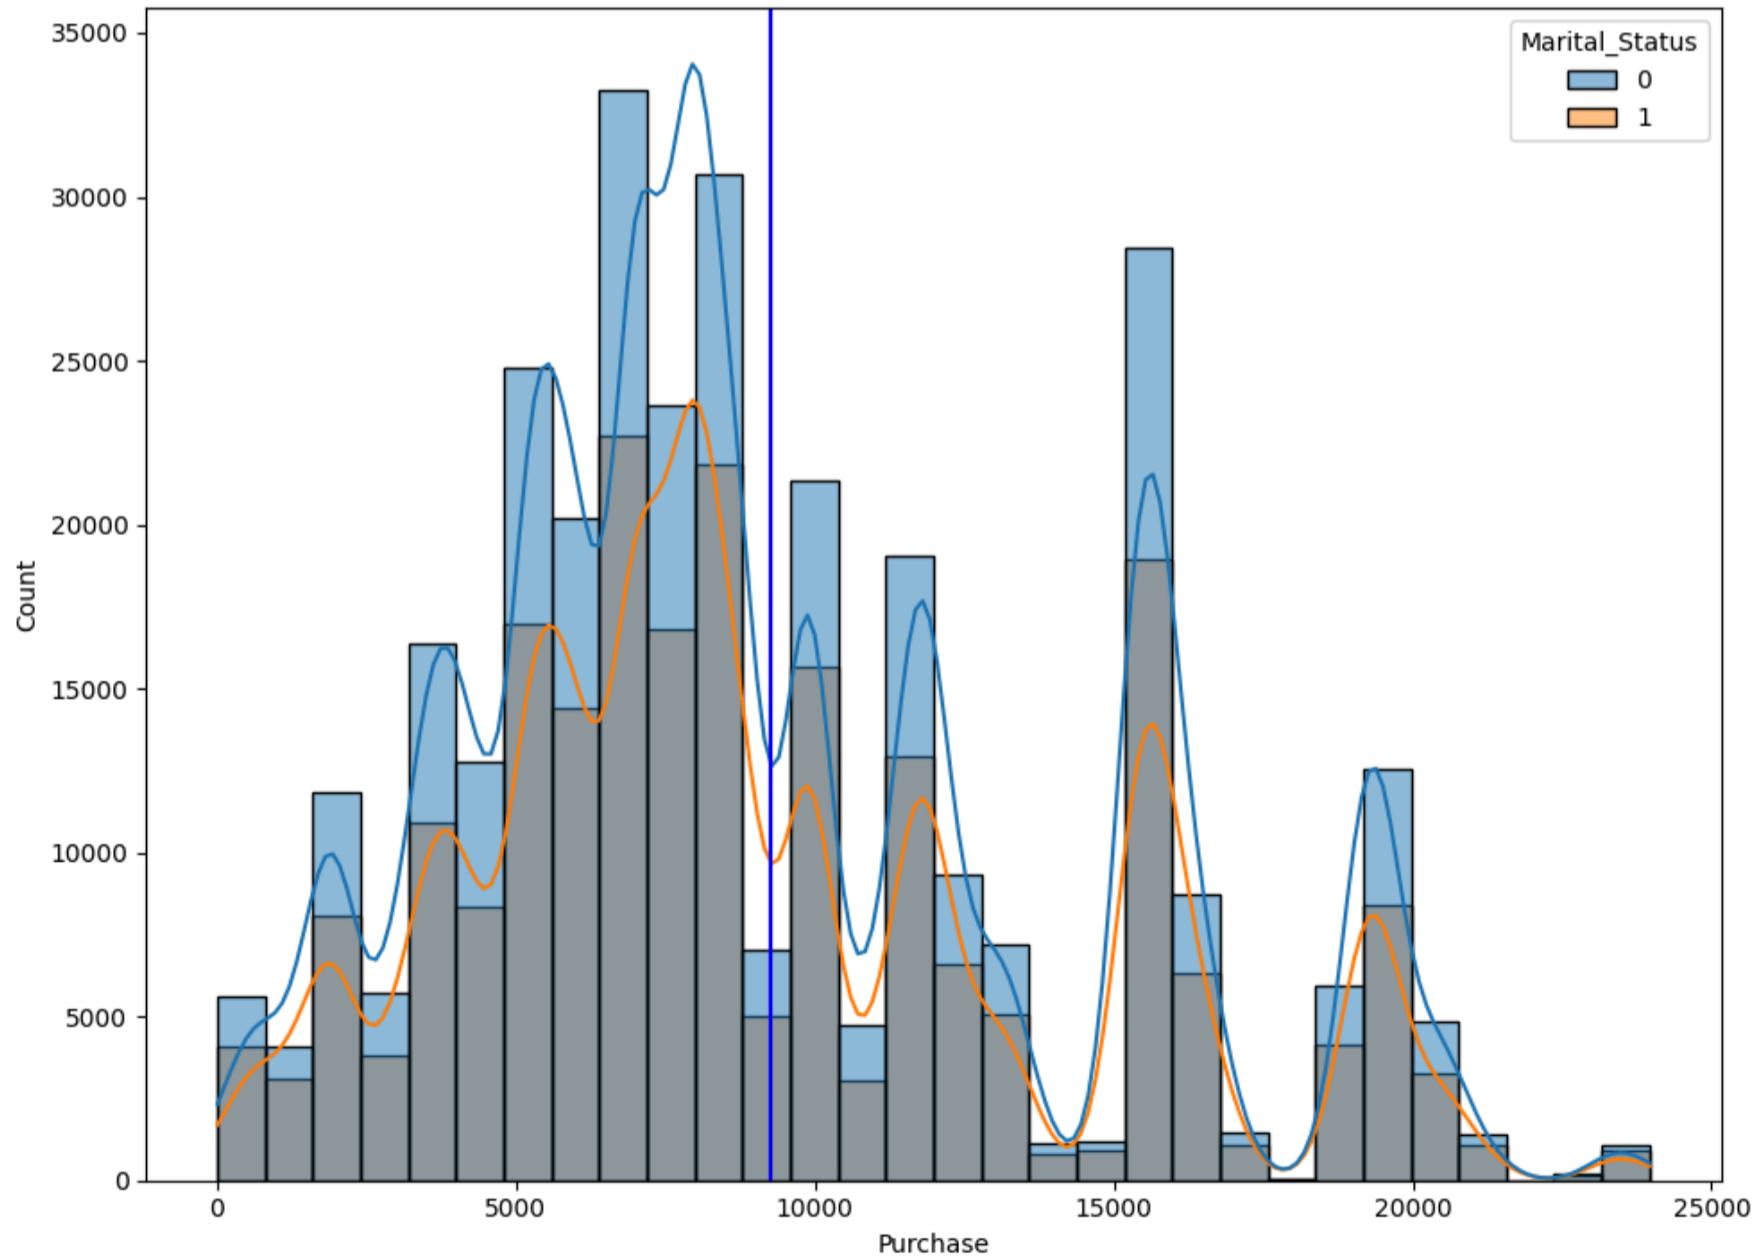
```
fig, axis = plt.subplots(1,1, figsize=(5,5))
sns.boxplot(data=df, x='Marital_Status', y='Purchase', ax=axis)
axis.set_title('MARITAL STATUS V/S PURCHASE DISTRIBUTION')
plt.tight_layout()
plt.show()
```

### MARITAL STATUS V/S PURCHASE DISTRIBUTION



```
fig, axis = plt.subplots(1,1, figsize=(10,7.5))
sns.histplot(data=df, x='Purchase', bins=30, kde=True, hue='Marital_Status', ax=axis)
plt.axvline(x=df['Purchase'].mean(), color='y')
plt.axvline(x=df[df['Marital_Status']==0]['Purchase'].mean(), color='r')
plt.axvline(x=df[df['Marital_Status']==1]['Purchase'].mean(), color='b')
plt.title('POPULATION MEAN : MARRIED V/S UNMARRIED')
plt.tight_layout()
plt.show()
```

### POPULATION MEAN : MARRIED V/S UNMARRIED



Notes: The distribution appears roughly normal, though outliers are evident around 15,000 and 20,000

### SAMPLE-DERIVED GENDER-CENTRIC ANALYSIS USING 1000 SAMPLE POINTS

#### MARRIED

```
df_married = df[df['Marital_Status']==1]
print("Mean of Purchase Amount by Married People:", df_married['Purchase'].mean())
print("Standard Deviation of Purchase Amount by Married People:", df_married['Purchase'].std())
```

```

sample_mean_of_purchase_by_married = []
for i in range(1000):
    sample_m = df_married['Purchase'].sample(n=7500)
    sample_m_mean = sample_m.mean()
    sample_mean_of_purchase_by_married.append(sample_m_mean)
sample_mean_of_purchase_by_married = np.array(sample_mean_of_purchase_by_married)

sample_mean_of_purchase_by_married[:5]

```

```

Mean of Purchase Amount by Married People: 9261.174574082374
Standard Deviation of Purchase Amount by Married People: 5016.89737779313
array([9307.33933333, 9249.242      , 9216.6016      , 9236.61186667,
       9303.73493333])

```

## UNMARRIED

```

df_unmarried = df[df['Marital_Status']==0]
print("Mean of Purchase Amount by Unmarried People:",df_unmarried['Purchase'].mean())
print("Standard Deviation of Purchase Amount by Unmarried People:",df_unmarried['Purchase'].std())

sample_mean_of_purchase_by_unmarried = []
for i in range(1000):
    sample_u = df_unmarried['Purchase'].sample(n=7500)
    sample_u_mean = sample_u.mean()
    sample_mean_of_purchase_by_unmarried.append(sample_u_mean)
sample_mean_of_purchase_by_unmarried = np.array(sample_mean_of_purchase_by_unmarried)

sample_mean_of_purchase_by_unmarried[:5]

```

```

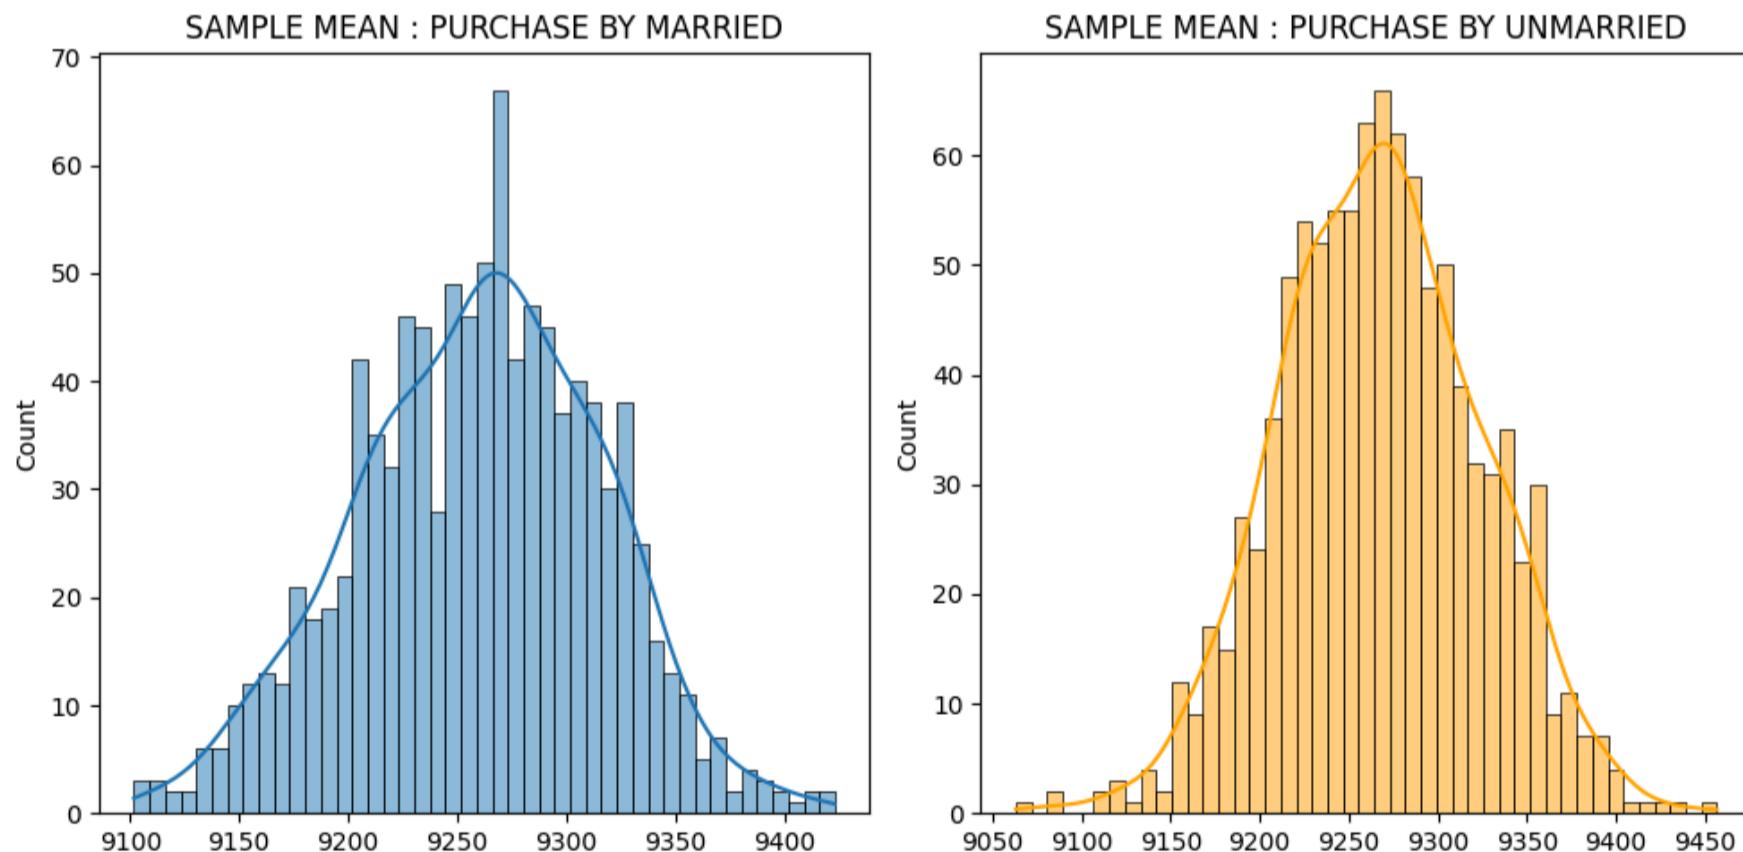
Mean of Purchase Amount by Unmarried People: 9265.907618921507
Standard Deviation of Purchase Amount by Unmarried People: 5027.347858674457
array([9259.70173333, 9303.49693333, 9296.08226667, 9255.4528      ,
       9273.15533333])

```

```

fig, axis = plt.subplots(1,2, figsize=(10,5))
sns.histplot(data=sample_mean_of_purchase_by_married, bins=45, kde=True, ax=axis[0])
axis[0].set_title('SAMPLE MEAN : PURCHASE BY MARRIED')
sns.histplot(data=sample_mean_of_purchase_by_unmarried, bins=45, kde=True, ax=axis[1], color='orange')
axis[1].set_title('SAMPLE MEAN : PURCHASE BY UNMARRIED')
plt.tight_layout()
plt.show()

```



## SAMPLE DISTRIBUTION MEANS:

```

print("Sample Mean of Purchase Amount by the Married:",sample_mean_of_purchase_by_married.mean())
print("Sample Mean of Purchase Amount by the Unmarried:",sample_mean_of_purchase_by_unmarried.mean())

```

```

Sample Mean of Purchase Amount by the Married: 9259.6269848
Sample Mean of Purchase Amount by the Unmarried: 9267.19743733334

```

The mean of this sampling distribution was found to be very close to the original population mean.

### Married

- Population mean = 9261.17
- Sampling distribution mean = 9259.62

### Unmarried

- Population mean = 9265.90

- Sampling distribution mean = 9267.19

This justifies the first property of the Central Limit Theorem: Sampling Distribution Mean  $\approx$  Population Mean

## ESTIMATION OF THE 95% CONFIDENCE INTERVAL

```
from scipy.stats import norm
```

### MARRIED

```
mu = df_married['Purchase'].mean()
sigma = df_married['Purchase'].std()
n = len(df_married)
alpha = 0.05
se = sigma/np.sqrt(n)
z1 = norm.ppf(alpha/2)
z2 = norm.ppf(1-alpha/2)
x1 = z1*se + mu
x2 = z2*se + mu
print("Confidence Interval for the Married:",(x1,x2))

Confidence Interval for the Married: (np.float64(9240.460427057078), np.float64(9281.888721107669))
```

### UNMARRIED

```
mu = df_unmarried['Purchase'].mean()
sigma = df_unmarried['Purchase'].std()
n = len(df_unmarried)
alpha = 0.05
se = sigma/np.sqrt(n)
z1 = norm.ppf(alpha/2)
z2 = norm.ppf(1-alpha/2)
x1 = z1*se + mu
x2 = z2*se + mu
print("Confidence Interval for the Unmarried:",(x1,x2))

Confidence Interval for the Unmarried: (np.float64(9248.61641818668), np.float64(9283.198819656332))
```

The **95% confidence intervals** in each case included the true population values — **consistent with the Central Limit Theorem**.

```
# H0 : Purchase amount is independent of Marital Status
# H1 : Purchase amount is dependent on Marital Status

sample_size = 10000

unmarried_group = df[df['Marital_Status'] == 0]['Purchase'].sample(n=sample_size, random_state=42)
married_group = df[df['Marital_Status'] == 1]['Purchase'].sample(n=sample_size, random_state=42)

result = ttest_ind(unmarried_group, married_group, alternative='two-sided')
alpha = 0.05

if result.pvalue < alpha:
    print("Reject Null Hypothesis")
    print("Marital Status has a significant impact on Purchase Amount\n")
else:
    print("Fail to Reject Null Hypothesis")
    print("Marital Status has no significant impact on Purchase Amount\n")

result

Fail to Reject Null Hypothesis
Marital Status has no significant impact on Purchase Amount

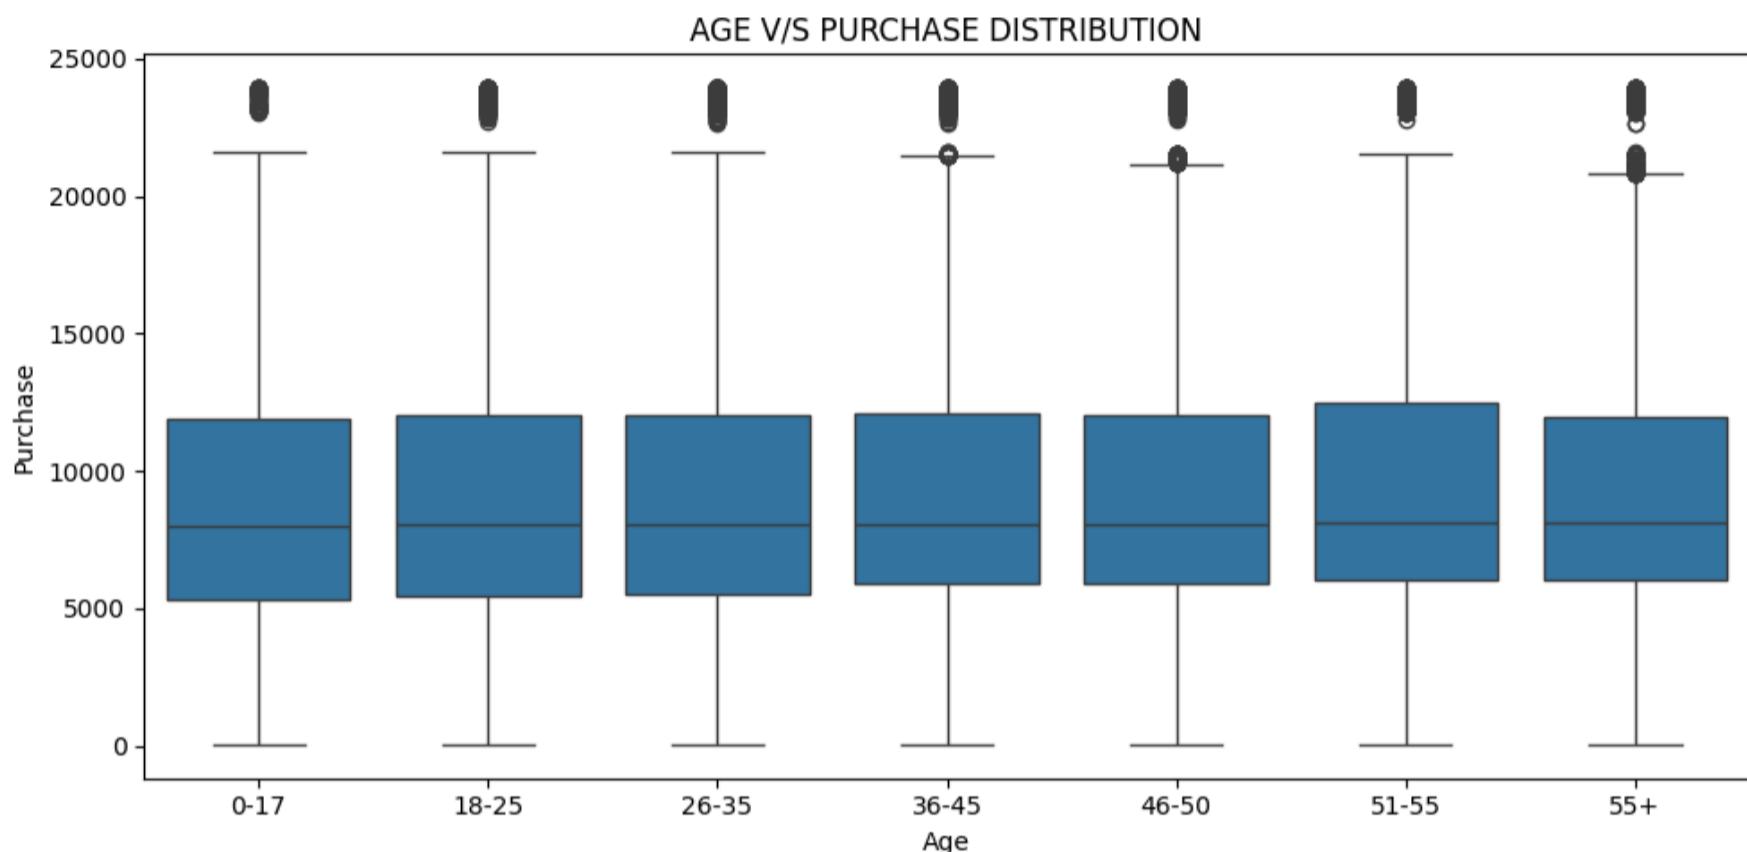
TtestResult(statistic=np.float64(0.5349013930718557), pvalue=np.float64(0.5927239922085221), df=np.float64(19998.0))
```

## ▼ Purchase vis-a-vis Age

```
df.groupby('Age')['Purchase'].describe()
```

	count	mean	std	min	25%	50%	75%	max
<b>Age</b>								
<b>0-17</b>	15102.0	8933.464640	5111.114046	12.0	5328.0	7986.0	11874.0	23955.0
<b>18-25</b>	99660.0	9169.663606	5034.321997	12.0	5415.0	8027.0	12028.0	23958.0
<b>26-35</b>	219587.0	9252.690633	5010.527303	12.0	5475.0	8030.0	12047.0	23961.0
<b>36-45</b>	110013.0	9331.350695	5022.923879	12.0	5876.0	8061.0	12107.0	23960.0
<b>46-50</b>	45701.0	9208.625697	4967.216367	12.0	5888.0	8036.0	11997.0	23960.0
<b>51-55</b>	38501.0	9534.808031	5087.368080	12.0	6017.0	8130.0	12462.0	23960.0
<b>55+</b>	21504.0	9336.280459	5011.493996	12.0	6018.0	8105.5	11932.0	23960.0

```
fig, axis = plt.subplots(1,1, figsize=(10,5))
sns.boxplot(data=df, x='Age', y='Purchase', ax=axis)
axis.set_title('AGE V/S PURCHASE DISTRIBUTION')
plt.tight_layout()
plt.show()
```



- The largest group is 26–35 years, followed by 18–25 years => Mid-career and young adults dominate the dataset
- Mean values rise with age up to the 51–55 group (9534.81, the highest) => Older adults tend to spend more on average
- Standard deviation is roughly the same across all groups (~5000)

```
from scipy.stats import f_oneway

# H0 : Purchase amount in independent of Age
# H1 : Purchase amount in dependent on Age

age_groups = [df[df['Age']==age_group]['Purchase'] for age_group in df['Age'].unique()]

result = f_oneway(*age_groups)
# The * before 'age_groups' in the function call unpacks the list of group arrays so that each group is passed as a separate argument
alpha = 0.05

if result.pvalue < alpha:
    print("Reject Null Hypothesis")
    print("Age has a significant impact on Purchase Amount\n")
else:
    print("Fail to Reject Null Hypothesis")
    print("Age has no significant impact on Purchase Amount\n")

result

Reject Null Hypothesis
Age has a significant impact on Purchase Amount

F_onewayResult(statistic=np.float64(40.57579909450407), pvalue=np.float64(1.053563939251671e-49))
```

## ▼ BUSINESS TAKEAWAY & ACTIONABLE INSIGHTS

- Target marketing and high-value promotions primarily toward **male and older adult segments** where spending is higher.
- Prioritize **City Category B** for inventory, ads, and store events due to its large customer share.

- **Marital status** can be **deprioritized** in segmentation, as its effect on purchase size is negligible.
- **Age-based marketing (especially for 36+ years)** may yield greater basket sizes and overall sales uplift.

These insights support more precise segmentation for Walmart's marketing and inventory planning, enhancing effectiveness during future Black Friday and major sales events.

Start coding or [generate](#) with AI.

Start coding or [generate](#) with AI.